

Q1. Query all columns for all American cities in the CITY table with populations larger than 100000.

The CountryCode for America is USA

City-Dataset:

<https://docs.google.com/spreadsheets/d/1dk9kRwcMxj5USuJqxtlTD05S-aOUD6fzNzVW41dcpgc/edit?usp=sharing>

The screenshot shows a database interface with a red border around the main content area. At the top, there are various icons for file operations, search, and navigation. To the right of the icons is a link to 'Limit to 1000 rows'. Below the toolbar, the SQL query is displayed:

```
1 • select * from city_table
2 where countrycode='USA' and population>100000
```

Below the query is a 'Result Grid' section. It includes a 'Result Grid' button, a 'Filter Rows:' button, and export options ('Export:' and 'Wrap'). The result grid itself has columns labeled 'ID', 'NAME', 'COUNTRYCODE', 'DISTRICT', and 'POPULATION'. The data is as follows:

	ID	NAME	COUNTRYCODE	DISTRICT	POPULATION
▶	3815	El Paso	USA	Texas	563662
	3878	Scottsdale	USA	Arizona	202705
	3965	Corona	USA	California	124966
	3973	Concord	USA	California	121780
	3977	Cedar Rapids	USA	Iowa	120758
	3982	Coral Springs	USA	Florida	117549

**Q2.** Query the NAME field for all American cities in the CITY table with populations larger than 120000. The CountryCode for America is USA.

The screenshot shows a MySQL Workbench interface. At the top, there is a toolbar with various icons for file operations, search, and navigation. To the right of the toolbar is a button labeled "Limit to 1000 rows". Below the toolbar, the SQL query is displayed:

```
1 • select name from city_table
2 where countrycode='USA' and population>120000
```

Below the query, there is a "Result Grid" section. It contains a table with one column labeled "name". The table displays the names of five cities: El Paso, Scottsdale, Corona, Concord, and Cedar Rapids. The "Scottsdale" row is currently selected, indicated by a blue background.

name
El Paso
Scottsdale
Corona
Concord
Cedar Rapids

**Q3.** Query all columns (attributes) for every row in the CITY table.

```
1 • select * from city_table
```

```
2
```

	ID	NAME	COUNTRYCODE	DISTRICT	POPULATION
▶	6	Rotterdam	NLD	Zuid-Holland	593321
	19	Zaanstad	NLD	Noord-Holland	135621
	214	Porto Alegre	BRA	Rio Grande do Sul	1314032
	397	Lauro de Freitas	BRA	Bahia	109236
	547	Dobric	BGR	Varna	100399
	552	Bujumbura	BDI	Bujumbura	300000
	554	Santiago de Chile	CHL	Santiago	4703954
	626	al-Minya	EGY	al-Minya	201360
	646	Santa Ana	SLV	Santa Ana	139389
	762	Bahir	Dar	ETH Amhara	96140
	796	Baguio	PHL	CAR	252386
	896	Malungon	PHL	Southern Mindanao	93232
	904	Banjul	GMB	Banjul	42326
	924	Villa	Nueva	GTM	101295
	...	...	...	...	...

Q4. Query all columns for a city in CITY with the ID 1661.

The screenshot shows a MySQL Workbench interface with a red border around the main content area. At the top, there's a toolbar with various icons: folder, table, lightning bolt, wrench, magnifying glass, hand, error, checkmark, close, and a refresh icon. To the right of the toolbar is a button labeled "Limit to 1000 rows". Below the toolbar, the SQL editor contains the following code:

```
1 • select * from city_table
2 where id=1661
3
```

Below the SQL editor is a horizontal line. Underneath it, there's a "Result Grid" section with the following interface elements:

- "Result Grid" tab
- Filter Rows: input field
- Export: icons for CSV, XML, and PDF
- Wi icon

The result grid itself has the following structure:

	ID	NAME	COUNTRYCODE	DISTRICT	POPULATION
▶	1661	Sayama	JPN	Saitama	162472

**Q5.** Query all attributes of every Japanese city in the CITY table. The COUNTRYCODE for Japan is JPN.

The screenshot shows a MySQL Workbench interface with a red border around the main content area. At the top, there's a toolbar with various icons: folder, file, lightning bolt, wrench, magnifying glass, hand, error, checkmark, close, and a gear icon. To the right of the icons is a dropdown menu set to "Limit to 1000 rows". Below the toolbar, a SQL editor window displays the following query:

```
1 • select * from city_table  
2 where countrycode='JPN'  
3
```

Below the SQL editor is a "Result Grid" section. It has a header row with columns: ID, NAME, COUNTRYCODE, DISTRICT, and POPULATION. The data grid contains five rows of results:

	ID	NAME	COUNTRYCODE	DISTRICT	POPULATION
▶	1613	Neyagawa	JPN	Osaka	257315
	1630	Ageo	JPN	Saitama	209442
	1661	Sayama	JPN	Saitama	162472
	1681	Omuta	JPN	Fukuoka	142889
	1739	Tokuyama	JPN	Yamaguchi	107078

Q6. Query the names of all the Japanese cities in the CITY table. The COUNTRYCODE for Japan is JPN

The screenshot shows a MySQL Workbench interface. At the top, there's a toolbar with various icons. Below the toolbar, a query editor window displays the following SQL code:

```
1 • select name from city_table
2 where countrycode='JPN'
3
```

The result grid below the query editor shows the following data:

	name
▶	Neyagawa
	Ageo
	Sayama
	Omura
	Tokuyama

**Station-table:**

<https://docs.google.com/spreadsheets/d/1sHPhE7walQD5mL7ppFNqybyoOJY3E51N0cWYzhp2UH4/edit?usp=sharing>

**Q7. Query a list of CITY and STATE from the STATION table. The STATION table is described as follows:**

where LAT\_N is the northern latitude and LONG\_W is the western longitude.

```
1 •      select city,state from station_data  
2
```

	city	state
▶	Kissee Mills	MO
	Loma Mar	CA
	Sandy Hook	CT
	Tipton	IN
	Arlington	CO
	Turner	AR
	Slidell	LA
	Negreet	LA
	Glencoe	KY
	Chelsea	IA
	Chignik Lag...	AK
	Pelahatchie	MS
	Hanna City	IL
	Dorrance	KS
•	All	All

Q8. Query a list of CITY names from STATION for cities that have an even ID number.  
Print the results in any order, but exclude duplicates from the answer.

The screenshot shows a MySQL Workbench interface. The top part is a query editor with a red border containing the following SQL code:

```
1 •  select distinct city
2      from station_data
3      where id%2=0
4      order by 1
5
```

The bottom part is a result grid titled "Result Grid" with a "Filter Rows:" button. It displays a list of city names:

	city
▶	Aguanga
	Alba
	Albany
	Amo
	Andersonville
	Archie
	Athens
	Atlantic Mine
	Bainbridge
	Baker
	Bass Harbor
	Bayville
	Beaufort
	Bellevue

At the bottom of the result grid, it says "station\_data 27 ×".

**Q9.** Find the difference between the total number of CITY entries in the table and the number of distinct CITY entries in the table.

The screenshot shows a MySQL Workbench interface. At the top, there is a toolbar with various icons for file operations, search, and navigation. Below the toolbar, a query window displays the following SQL code:

```
1 • select count(city)-count(distinct city) as diff
2   from station_data
3
```

The code uses a self-explanatory alias 'diff' for the result of the subtraction. The 'distinct' keyword ensures that each city name is counted only once, even if it appears multiple times in the 'station\_data' table.

Below the query window, there is a results grid labeled 'Result Grid'. The grid has one row and two columns. The first column is labeled 'diff' and contains the value '13'. There are also buttons for 'Filter Rows' and 'Export'.

**Q10.** Query the two cities in STATION with the shortest and longest CITY names, as well as their respective lengths (i.e.: number of characters in the name). If there is more than one smallest or largest city, choose the one that comes first when ordered alphabetically

Query 1 x

```
1 select city,length from (
2     select city ,length(city) as length,dense_rank() over (order by length(city)) as min_to_max ,dense_rank() over (order by length(city) desc)
3         as max_to_min
4     from station_data
5 )b
6 where min_to_max =1 or max_to_min=1
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

city	length
White Horse Beach	18
West Baden Springs	18
Roy	3
Amo	3
Lee	3

Result Grid  
Form Editor

**Q11. Query the list of CITY names starting with vowels (i.e., a, e, i, o, or u) from STATION.  
Your result cannot contain duplicates.**

```
1
2 •   SELECT DISTINCT city
3     FROM station_data
4    WHERE city RLIKE '^[aeiouAEIOU].*[aeiouAEIOU]$'
5    order by 1
```

	city
▶	Acme
	Aguanga
	Alba
	Aliso Viejo
	Alpine
	Amazonia
	Amo
	Andersonville
	Archie
	Arispe
	Arkadelphia
	Atlantic Mine
	East China
	East Irvine
	Endeavor

**Q12. Query the list of CITY names ending with vowels (a, e, i, o, u) from STATION. Your result cannot contain duplicates.**



```
1 •  SELECT DISTINCT city
2      FROM station_data
3      WHERE city RLIKE '[aeiouAEIOU]$'
4      order by 1
```

Result Grid



Filter Rows:

Export:

	city
▶	Acme
	Aguanga
	Alba
	Aliso Viejo
	Alpine
	Amazonia
	Amo
	Andersonville
	Archie
	Arispe
	Arkadelphia
	Atlantic Mine
	Baileyville
	Bainbridge
	... 51 more
station_data 51 ×	

Q13. Query the list of CITY names from STATION that do not start with vowels. Your result cannot contain duplicates.

The screenshot shows a MySQL Workbench interface with a red border around the main content area. At the top, there is a toolbar with various icons. Below the toolbar, a code editor window displays the following SQL query:

```
1 •  SELECT DISTINCT city
2   FROM station_data
3 WHERE city NOT RLIKE '^[aeiouAEIOU]'
4   ORDER BY 1
5
6   -- .*[aeiouAEIOU]$
```

Below the code editor is a results grid titled "Result Grid". The grid has a single column labeled "city". The data in the grid is as follows:

city
Baileyville
Bainbridge
Baker
Baldwin
Barrigada
Bass Harbor
Baton Rouge
Bayville
Beaufort
Beaver Island
Bellevue
Benedict
Bennington
Bentonville

At the bottom of the results grid, it says "station\_data 52 ×". Above the results grid, there are buttons for "Result Grid", "Filter Rows:", "Export:", and "Wrap Cell Content".

Q14. Query the list of CITY names from STATION that do not end with vowels. Your result cannot contain duplicates.

Query 1   |     |   | Limit to 100

```
1 •   SELECT DISTINCT city
2     FROM station_data
3     WHERE city NOT RLIKE '[aeiouAEIOU]$'
4     ORDER BY 1
5
```

Result Grid 

Filter Rows:

Export: 

	city
▶	Addison
	Agency
	Alanson
	Albany
	Albion
	Algonac
	Allerton
	Alton
	Andover
	Anthony
	Arlington
	Arrows...
	Athens
	Auburn
	Belo...

station\_data 54 

Q15. Query the list of CITY names from STATION that either do not start with vowels or do not end with vowels. Your result cannot contain duplicates.

Query 1 x

File Edit View Insert Cell Tools Help

1 • SELECT DISTINCT city FROM station\_data WHERE  
2 city RLIKE '^[^aeiouAEIOU].\*|.\*[^AEIOUaeiou]\$';  
3

Result Grid | Filter Rows: Export: Wrap Cell Co

city
Kissee Mills
Loma Mar
Sandy Hook
Tipton
Arlington
Turner
Slidell
Negreet
Glencoe
Chelsea
Chignik Lag...
Pelahatchie
Hanna City
Dorrance

station\_data 55 x

**Q16. Query the list of CITY names from STATION that do not start with vowels and do not end with vowels. Your result cannot contain duplicates.**

Query 1 ×

1 • SELECT DISTINCT city  
2 FROM station\_data  
3 WHERE city NOT LIKE 'A%' AND city NOT LIKE 'E%' AND city NOT LIKE 'I%' AND city NOT LIKE 'O%' AND city NOT LIKE 'U%'  
4 AND city NOT LIKE '%A' AND city NOT LIKE '%E' AND city NOT LIKE '%I' AND city NOT LIKE '%O' AND city NOT LIKE '%U'  
5 order by 1  
6

Result Grid | Filter Rows: \_\_\_\_\_ | Export: | Wrap Cell Content:

city
Baker
Baldwin
Bass Harbor
Beaufort
Beaver Island
Benedict
Bennington
Berryton
Beverly
Bison
Blue River
Bowdon
Bowdon Junc...
Bridgeport
Bridgeport

station\_data 75 ×

**Q17.**

Table: Product

Column Name	Type
product_id	int
product_name	varchar
unit_price	int

product\_id is the primary key of this table.

Each row of this table indicates the name and the price of each product.

Table: Sales

Column Name	Type
seller_id	int
product_id	int
buyer_id	int
sale_date	date
quantity	int
price	int

This table has no primary key, it can have repeated rows.

product\_id is a foreign key to the Product table.  
Each row of this table contains some information about one sale.

Write an SQL query that reports the products that were only sold in the first quarter of 2019. That is, between 2019-01-01 and 2019-03-31 inclusive.

Return the result table in any order.

The query result format is in the following example.

Input:

Product table:

product_id	product_name	unit_price
1	S8	1000
2	G4	800
3	iPhone	1400

Sales table:

seller_id	product_id	buyer_id	sale_date	quantity	price
1	1	1	2019-01-21	2	2000
1	2	2	2019-02-17	1	800
2	2	3	2019-06-02	1	800
3	3	4	2019-05-13	2	2800

Output:

product_id	product_name
1	S8

Explanation:

The product with id 1 was only sold in the spring of 2019.

The product with id 2 was sold in the spring of 2019 but was also sold after the spring of 2019.

The product with id 3 was sold after spring 2019.

We return only product 1 as it is the product that was only sold in the spring of 2019.

Query 1 x

14  
15 • select product.product\_id,product\_name,MAX(sale\_date) as latest\_sale  
16 from sales  
17 join product on product.product\_id=sales.product\_id  
18 group by 1,2  
19 having latest\_sale between '2019-01-01' and '2019-03-31'  
20  
21  
22

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

	product_id	product_name	latest_sale
▶	1	S8	2019-01-21

**Q18.**

Table: Views

Column Name	Type
article_id	int
author_id	int
viewer_id	int
view_date	date

There is no primary key for this table, it may have duplicate rows.

Each row of this table indicates that some viewer viewed an article (written by some author) on some date.

Note that equal author\_id and viewer\_id indicate the same person.

Write an SQL query to find all the authors that viewed at least one of their own articles.

Return the result table sorted by id in ascending order.

The query result format is in the following example.

Input:

Views table:

article_id	author_id	viewer_id	view_date
1	3	5	2019-08-01
1	3	6	2019-08-02
2	7	7	2019-08-01
2	7	6	2019-08-02
4	7	1	2019-07-22
3	4	4	2019-07-21
3	4	4	2019-07-21

Output:

id
4
7

```
18
19 •   select distinct views.author_id from views
20     join views as v on v.viewer_id=views.author_id
21     order by 1 asc
22
23
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	author_id			
▶	4			
	7			

### Q19.

Table: Delivery

Column Name	Type
delivery_id	int
customer_id	int
order_date	date
customer_pref_delivery_date	date

`delivery_id` is the primary key of this table.

The table holds information about food delivery to customers that make orders at some date and specify a preferred delivery date (on the same order date or after it).

If the customer's preferred delivery date is the same as the order date, then the order is called immediately; otherwise, it is called scheduled.

Write an SQL query to find the percentage of immediate orders in the table, rounded to 2 decimal places.

The query result format is in the following example.

Input:

Delivery table:

<code>delivery_id</code>	<code>customer_id</code>	<code>order_date</code>	<code>customer_pref_delivery_date</code>
1	1	2019-08-01	2019-08-02
2	5	2019-08-02	2019-08-02
3	1	2019-08-11	2019-08-11
4	3	2019-08-24	2019-08-26
5	4	2019-08-21	2019-08-22
6	2	2019-08-11	2019-08-13

Output:

<code>immediate_percentage</code>
33.33

Explanation: The orders with delivery id 2 and 3 are immediate while the others are scheduled.

The screenshot shows a MySQL Workbench interface with a red border around the main content area. At the top, there's a toolbar with various icons. Below the toolbar, the SQL editor contains the following code:

```

21
22
23 • select round(count(case when order_date=customer_pref_delivery_date then 1 end)/count(*),2) as percent_immediate
24   from delivery
25
26
27
28
29

```

Below the code, the results are displayed in a grid:

percent_immediate
0.33

## Q20.

Table: Ads

Column Name	Type
ad_id	int
user_id	int
action	enum

(ad\_id, user\_id) is the primary key for this table.

Each row of this table contains the ID of an Ad, the ID of a user, and the action taken by this user regarding this Ad.

The action column is an ENUM type of ('Clicked', 'Viewed', 'Ignored').

A company is running Ads and wants to calculate the performance of each Ad.

Performance of the Ad is measured using Click-Through Rate (CTR) where:

$$CTR = \begin{cases} 0, & \text{if Ad total clicks + Ad total views} = 0 \\ \frac{\text{Ad total clicks}}{\text{Ad total clicks} + \text{Ad total views}} \times 100, & \text{otherwise} \end{cases}$$

Write an SQL query to find the ctr of each Ad. Round ctr to two decimal points.

Return the result table ordered by ctr in descending order and by ad\_id in ascending order in case of a tie.

The query result format is in the following example.

Input:

Ads table:

ad_id	user_id	action
1	1	Clicked
2	2	Clicked
3	3	Viewed
5	5	Ignored
1	7	Ignored
2	7	Viewed
3	5	Clicked
1	4	Viewed
2	11	Viewed
1	2	Clicked

Output:

ad_id	ctr
1	66.67
3	50
2	33.33
5	0

Explanation:

for ad\_id = 1, ctr =  $(2/(2+1)) * 100 = 66.67$

for ad\_id = 2, ctr =  $(1/(1+2)) * 100 = 33.33$

for ad\_id = 3, ctr =  $(1/(1+1)) * 100 = 50.00$

for ad\_id = 5, ctr = 0.00, Note that ad\_id = 5 has no clicks or views.

Note that we do not care about Ignored Ads.

```
52 •   select ad_id,
53     case
54     when (count(case when action='Viewed' then 1 end)+count(case when action='Clicked' then 1 end))=0 then 0
55     else count(case when action='Clicked' then 1 end)/(count(case when action='Viewed' then 1 end)+count(case when action='Clicked' then 1 end
56     ))*100 end as ctr
57   from ads
58   group by 1
59
60
--
```

result Grid | Filter Rows: | Export: | Wrap Cell Content: |  | Result Grid | Form Editor

ad_id	ctr
1	66.6667
2	33.3333
3	50.0000
5	0

**Q21.**

Table: Employee

Column Name	Type
employee_id	int
team_id	int

employee\_id is the primary key for this table.

Each row of this table contains the ID of each employee and their respective team.

Write an SQL query to find the team size of each of the employees.

Return result table in any order.

The query result format is in the following example.

Input:

Employee Table:

employee_id	team_id
1	8
2	8
3	8
4	7
5	9
6	9

Output:

employee_id	team_size
1	3
2	3
3	3
4	1
5	2
6	2

Explanation:

Employees with Id 1,2,3 are part of a team with team\_id = 8.

An employee with Id 4 is part of a team with team\_id = 7.

Employees with Id 5,6 are part of a team with team\_id = 9.

```
77 •   select employee_id, count(*) over(partition by team_id) as team_size  
78     from employee  
79    order by 1  
80  
81  
82  
83
```

Result Grid | Filter Rows:  Export: Wrap Cell Content:

	employee_id	team_size
▶	1	3
	2	3
	3	3
	4	1
	5	2
	6	2

**Q22.**

Table: Countries

Column Name	Type
country_id	int
country_name	varchar

country\_id is the primary key for this table.

Each row of this table contains the ID and the name of one country.

Table: Weather

Column Name	Type
country_id	int
weather_state	int
day	date

(country\_id, day) is the primary key for this table.

Each row of this table indicates the weather state in a country for one day.

Write an SQL query to find the type of weather in each country for November 2019.

The type of weather is:

- Cold if the average weather\_state is less than or equal 15,
- Hot if the average weather\_state is greater than or equal to 25, and
- Warm otherwise.

Return result table in any order.

The query result format is in the following example.

Input:

Countries table:

country_id	country_name
2	USA
3	Australia
7	Peru
5	China
8	Morocco
9	Spain

Weather table:

country_id	weather_state	day
2	15	2019-11-01
2	12	2019-10-28

2	12	2019-10-27
3	-2	2019-11-10
3	0	2019-11-11
3	3	2019-11-12
5	16	2019-11-07
5	18	2019-11-09
5	21	2019-11-23
7	25	2019-11-28
7	22	2019-12-01
7	20	2019-12-02
8	25	2019-11-05
8	27	2019-11-15
8	31	2019-11-25
9	7	2019-10-23
9	3	2019-12-23

Output:

country_name	weather_type
USA	Cold
Australia	Cold
Peru	Hot
Morocco	Hot
China	Warm

Explanation:

Average weather\_state in the USA in November is  $(15) / 1 = 15$  so the weather type is Cold.

Average weather\_state in Australia in November is  $(-2 + 0 + 3) / 3 = 0.333$  so the weather type is Cold.

Average weather\_state in Peru in November is  $(25) / 1 = 25$  so the weather type is Hot.

The average weather\_state in China in November is  $(16 + 18 + 21) / 3 = 18.333$  so the weather type is warm.

Average weather\_state in Morocco in November is  $(25 + 27 + 31) / 3 = 27.667$  so the weather type is Hot.

We know nothing about the average weather\_state in Spain in November so we do not include it in the result table.

```
125 •   select country_name,  
126     case  
127       when avg(weather_state)>=25 then 'Hot'  
128       when avg(weather_state)<=15 then 'Cold'  
129       else 'Warm' end as weather_type  
130   from weather  
131   join countries on countries.country_id=weather.country_id  
132   where extract(month from day)=11 and extract(year from day)=2019  
133   group by 1  
134
```

Result Grid | Filter Rows: \_\_\_\_\_ | Export: | Wrap Cell Content:

country_name	weather_type
USA	Cold
AUSTRALIA	Cold
CHINA	Warm
PERU	Hot
MOROCCO	Hot

**Q23.**

Table: Prices

Column Name	Type
product_id	int
start_date	date
end_date	date
price	int

(product\_id, start\_date, end\_date) is the primary key for this table.

Each row of this table indicates the price of the product\_id in the period from start\_date to end\_date.  
For each product\_id there will be no two overlapping periods. That means there will be no two intersecting periods for the same product\_id.

Table: UnitsSold

Column Name	Type
product_id	int
purchase_date	date
units	int

There is no primary key for this table, it may contain duplicates.

Each row of this table indicates the date, units, and product\_id of each product sold.

Write an SQL query to find the average selling price for each product. average\_price should be rounded to 2 decimal places.

Return the result table in any order.

The query result format is in the following example.

**Input:**

Prices table:

product_id	start_date	end_date	price
1	2019-02-17	2019-02-28	5
1	2019-03-01	2019-03-22	20
2	2019-02-01	2019-02-20	15
2	2019-02-21	2019-03-31	30

UnitsSold table:

product_id	purchase_date	units
1	2019-02-25	100
1	2019-03-01	15
2	2019-02-10	200
2	2019-03-22	30

**Output:**

product_id	average_price
1	6.96
2	16.96

**Explanation:**

Average selling price = Total Price of Product / Number of products sold.

Average selling price for product 1 =  $((100 * 5) + (15 * 20)) / 115 = 6.96$

Average selling price for product 2 =  $((200 * 15) + (30 * 30)) / 230 = 16.96$

```
171 • select prices.product_id,round(sum(price*units)/sum(units),2) as sold_price from units_sold  
172     join prices on prices.product_id=units_sold .product_id and purchase_date between start_date and end_date  
173     group by 1  
174  
175  
176 |  
177  
178  
179
```

Result Grid | Filter Rows:  Export: Wrap Cell Content:

	product_id	sold_price
▶	1	6.96
	2	16.96

**Q25.**

Table: Activity

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player\_id, event\_date) is the primary key of this table.

This table shows the activity of players of some games.

Each row is a record of a player who logged in and played a number of games (possibly 0) before logging out on someday using some device.

Write an SQL query to report the device that is first logged in for each player.

Return the result table in any order.

The query result format is in the following example.

Input:

Activity table:

player_id	device_id	event_date	games_played
1	2	2016-03-01	5
1	2	2016-05-02	6
2	3	2017-06-25	1
3	1	2016-03-02	0
3	4	2018-07-03	5

Output:

player_id	device_id
1	2
2	3
3	1

```
192      -- 25
193 •  select player_id,min(event_date) as first_login
194    from activity
195   group by 1
196
197
198
```

Result Grid



Filter Rows:

Export:



Wrap Cell Co

	player_id	first_login
▶	1	2016-03-01
	2	2017-06-25
	3	2016-03-02

**Q26.**

Table: Products

Column Name	Type
product_id	int
product_name	varchar
product_category	varchar

product\_id is the primary key for this table.

This table contains data about the company's products.

Table: Orders

Column Name	Type
product_id	int
order_date	date
unit	int

There is no primary key for this table. It may have duplicate rows.

product\_id is a foreign key to the Products table.

unit is the number of products ordered in order\_date.

Write an SQL query to get the names of products that have at least 100 units ordered in February 2020 and their amount.

Return result table in any order.

The query result format is in the following example.

Input:

Products table:

product_id	product_name	product_category
1	Leetcode Solutions	Book
2	Jewels of Stringology	Book
3	HP	Laptop
4	Lenovo	Laptop
5	Leetcode Kit	T-shirt

Orders table:

product_id	order_date	unit
1	2020-02-05	60
1	2020-02-10	70
2	2020-01-18	30
2	2020-02-11	80
3	2020-02-17	2
3	2020-02-24	3
4	2020-03-01	20
4	2020-03-04	30
4	2020-03-04	60
5	2020-02-25	50
5	2020-02-27	50
5	2020-03-01	50

Output:

product_name	unit
Leetcode Solutions	130
Leetcode Kit	100

Explanation:

Products with product\_id = 1 is ordered in February a total of  $(60 + 70) = 130$ .

Products with product\_id = 2 is ordered in February a total of 80.

Products with product\_id = 3 is ordered in February a total of  $(2 + 3) = 5$ .

Products with product\_id = 4 was not ordered in February 2020.

Products with product\_id = 5 is ordered in February a total of  $(50 + 50) = 100$ .

- - -

```

238      -- Ans 27
239 •   select orders.product_id,product_name,product_category,sum(unit) as units
240     from products
241     join orders on orders.product_id=products.product_id
242     and extract(month from order_date)=2
243     group by 1,2,3
244     having sum(unit)>=100

```

Result Grid | Filter Rows:  Export: Wrap Cell Content:

	product_id	product_name	product_category	units
▶	1	LEETCODE SOLUTIONS	BOOK	130
	5	LEETCODE KIT	T-SHIRT	100

## Q27.

Table: Users

Column Name	Type
user_id	int
name	varchar
mail	varchar

user\_id is the primary key for this table.

This table contains information of the users signed up in a website. Some emails are invalid.

Write an SQL query to find the users who have valid emails.

A valid e-mail has a prefix name and a domain where:

- The prefix name is a string that may contain letters (upper or lower case), digits, underscore '\_', period '.', and/or dash '-'. The prefix name must start with a letter.
- The domain is '@leetcode.com'.

Return the result table in any order.

The query result format is in the following example.

Input:

Users table:

user_id	name	mail
1	Winston	winston@leetcode.com
2	Jonathan	jonathanisgreat
3	Annabelle	bella-@leetcode.com
4	Sally	sally.come@leetcode.com
5	Marwan	quarz#2020@leetcode.com
6	David	david69@gmail.com
7	Shapiro	.shapo@leetcode.com

Output:

user_id	name	mail
1	Winston	winston@leetcode.com
3	Annabelle	bella-@leetcode.com
4	Sally	sally.come@leetcode.com

Explanation:

The mail of user 2 does not have a domain.

The mail of user 5 has the # sign which is not allowed.

The mail of user 6 does not have the leetcode domain.

The mail of user 7 starts with a period.

```
264      -- Ans 27
265 •  select *
266   from users
267   where mail regexp '^[A-Za-z][A-Za-z0-9\_.\_-]*@leetcode\.com$'
268
269
270
271
272
273
```

Result Grid			
	user_id	name	mail
▶	1	WINSTON	winston@leetcode.com
	3	ANNABELLE	bella-@leetcode.com
	4	SALLY	sally.come@leetcode.com
*	NULL	NULL	NULL

**Q28.**

Table: Customers

Column Name	Type
customer_id	int
name	varchar
country	varchar

customer\_id is the primary key for this table.

This table contains information about the customers in the company.

Table: Product

Column Name	Type
customer_id	int
name	varchar
country	varchar

product\_id is the primary key for this table.  
 This table contains information on the products in the company.  
 price is the product cost.

Table: Orders

Column Name	Type
order_id	int
customer_id	int
product_id	int
order_date	date
quantity	int

order\_id is the primary key for this table.  
 This table contains information on customer orders.  
 customer\_id is the id of the customer who bought "quantity" products with id "product\_id".  
 Order\_date is the date in format ('YYYY-MM-DD') when the order was shipped.

Write an SQL query to report the customer\_id and customer\_name of customers who have spent at least \$100 in each month of June and July 2020.  
 Return the result table in any order.  
 The query result format is in the following example.

Input:

Customers table:

customer_id	name	country
1	Winston	USA
2	Jonathan	Peru
3	Moustafa	Egypt

Product table:

product_id	description	price
10	LC Phone	300
20	LC T-Shirt	10
30	LC Book	45
40	LC Keychain	2

Orders table:

order_id	customer_id	product_id	order_date	quantity
1	1	10	2020-06-10	1
2	1	20	2020-07-01	1
3	1	30	2020-07-08	2
4	2	10	2020-06-15	2
5	2	40	2020-07-01	10
6	3	20	2020-06-24	2
7	3	30	2020-06-25	2
9	3	30	2020-05-08	3

Output:

customer_id	name
1	Winston

Explanation:

Winston spent \$300 ( $300 * 1$ ) in June and \$100 ( $10 * 1 + 45 * 2$ ) in July 2020.

Jonathan spent \$600 ( $300 * 2$ ) in June and \$20 ( $2 * 10$ ) in July 2020.

Moustafa spent \$110 ( $10 * 2 + 45 * 2$ ) in June and \$0 in July 2020.

## Ans 28

```
select customer_id, name from (
    select customers.customer_id, name, extract(month from
        order_date), sum(quantity*price) as amount_spent, row_number() over(partition
        by customers.customer_id) as instances
    from orders
    join customers on customers.customer_id=orders.customer_id
    join products on products.product_id=orders.product_id
    and extract(month from order_date) in (6,7)
    group by 1,2,3
    having amount_spent>=100
)b
where instances=2
```

**Q29.**

Table: TVProgram

Column Name	Type
program_date	date
content_id	int
channel	varchar

(program\_date, content\_id) is the primary key for this table.

This table contains information about the programs on the TV.

content\_id is the id of the program in some channel on the TV.

Table: Content

Column Name	Type
content_id	varchar
title	varchar
Kids_content	enum
content_type	varchar

content\_id is the primary key for this table.

Kids\_content is an enum that takes one of the values ('Y', 'N') where:

'Y' means content for kids, otherwise 'N' is not content for kids.  
content\_type is the category of the content as movies, series, etc.

Write an SQL query to report the distinct titles of the kid-friendly movies streamed in June 2020.

Return the result table in any order.

The query result format is in the following example.

Input:

TVProgram table:

program_date	content_id	channel
2020-06-10 08:00	1	LC-Channel
2020-05-11 12:00	2	LC-Channel
2020-05-12 12:00	3	LC-Channel
2020-05-13 14:00	4	Disney Ch
2020-06-18 14:00	4	Disney Ch
2020-07-15 16:00	5	Disney Ch

Content table:

content_id	title	Kids_content	content_type
1	Leetcode Movie	N	Movies
2	Alg. for Kids	Y	Series
3	Database Sols	N	Series
4	Aladdin	Y	Movies
5	Cinderella	Y	Movies

Output:

title
Aladdin

Explanation:

"Leetcode Movie" is not a content for kids.

"Alg. for Kids" is not a movie.

"Database Sols" is not a movie

"Aladdin" is a movie, content for kids and was streamed in June 2020.

"Cinderella" was not streamed in June 2020.

```
403    -- Ans 29
404 • select distinct title from tv_program
405   join content on content.content_id=tv_program.content_id
406   and extract(month from program_date)=6 and extract(year from program_date)=2020
407   and kids_content='Y'
408
409
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	title			
▶	ALADDIN			

**Q30.**

Table: NPV

Column Name	Type
id	int
year	int
npv	int

(id, year) is the primary key of this table.

The table has information about the id and the year of each inventory and the corresponding net present value.

Table: Queries

Column Name	Type
id	int
year	int

(id, year) is the primary key of this table.

The table has information about the id and the year of each inventory query.

Write an SQL query to find the npv of each query of the Queries table.

Return the result table in any order.

The query result format is in the following example.

Input:

NPV table:

id	year	npv
1	2018	100
7	2020	30
13	2019	40
1	2019	113
2	2008	121
3	2009	12
11	2020	99
7	2019	0

Queries table:

id	year
1	2019
2	2008
3	2009
7	2018
7	2019
7	2020
13	2019

Output:

id	year	npv
1	2019	113
2	2008	121
3	2009	12
7	2018	0
7	2019	0
7	2020	30
13	2019	40

Explanation:

The npv value of (7, 2018) is not present in the NPV table, we consider it 0.

The npv values of all other queries can be found in the NPV table.

Query 1 x

449 • select queries.id,queries.year,npv from queries  
450 left Join npv on npv.year=queries.year and npv.id=queries.id  
451  
452  
453  
454

Result Grid | Filter Rows: Export: Wrap Cell Content:

	id	year	npv
▶	1	2019	113
	2	2008	121
	3	2009	12
	7	2018	NULL
	7	2019	0
	7	2020	30
	13	2019	40

**Q31.**

Table: NPV

Column Name	Type
id	int
year	int
npv	int

(id, year) is the primary key of this table.

The table has information about the id and the year of each inventory and the corresponding net present value.

Table: Queries

Column Name	Type
id	int
year	int

(id, year) is the primary key of this table.

The table has information about the id and the year of each inventory query.

Write an SQL query to find the npv of each query of the Queries table.

Return the result table in any order.

The query result format is in the following example.

Input:

NPV table:

id	year	npv
1	2018	100
7	2020	30
13	2019	40
1	2019	113
2	2008	121
3	2009	12
11	2020	99
7	2019	0

Queries table:

id	year
1	2019
2	2008
3	2009
7	2018
7	2019
7	2020
13	2019

Output:

id	year	npv
1	2019	113
2	2008	121
3	2009	12
7	2018	0
7	2019	0
7	2020	30
13	2019	40

```
449 •  select queries.id,queries.year,npv from queries
450      left Join npv on npv.year=queries.year and npv.id=queries.id
451
452
453
454
```

Result Grid | Filter Rows:  Export: Wrap Cell Content:

	id	year	npv
▶	1	2019	113
	2	2008	121
	3	2009	12
	7	2018	NULL
	7	2019	0
	7	2020	30
	13	2019	40

**Q32.**

Table: Employees

Column Name	Type
id	int
name	varchar

id is the primary key for this table.

Each row of this table contains the id and the name of an employee in a company.

Table: EmployeeUNI

Column Name	Type
id	int
unique_id	int

(id, unique\_id) is the primary key for this table.

Each row of this table contains the id and the corresponding unique id of an employee in the company.

Write an SQL query to show the unique ID of each user, If a user does not have a unique ID replace just show null.

Return the result table in any order.

The query result format is in the following example.

Input:

Employees table:

id	name
1	Alice
7	Bob
11	Meir
90	Winston
3	Jonathan

EmployeeUNI table:

id	unique_id
3	1
11	2
90	3

Output:

unique_id	name
null	Alice
null	Bob
2	Meir
3	Winston
1	Jonathan

Explanation:

Alice and Bob do not have a unique ID, We will show null instead.

The unique ID of Meir is 2.

The unique ID of Winston is 3.

The unique ID of Jonathan is 1.

475

```
476 •   select employees.id, name, unique_id from employees  
477     left join employees_uni on employees_uni.id=employees.id  
478  
479  
480
```

Result Grid | Filter Rows:  Export: Wrap Cell Content:

	id	name	unique_id
▶	1	ALICE	NULL
	3	JONATHAN	1
	7	BOB	NULL
	11	MEIR	2
	90	WINSTON	3

**Q33.**

Table: Users

Column Name	Type
id	int
name	varchar

id is the primary key for this table.

name is the name of the user.

Table: Rides

Column Name	Type
id	int
user_id	int
distance	int

id is the primary key for this table.

user\_id is the id of the user who travelled the distance "distance".

Write an SQL query to report the distance travelled by each user.

Return the result table ordered by travelled\_distance in descending order, if two or more users travelled the same distance, order them by their name in ascending order.

The query result format is in the following example.

Input:

Users table:

id	name
1	Alice
2	Bob
3	Alex
4	Donald
7	Lee

13	Jonathan
19	Elvis

Rides table:

id	user_id	distance
1	1	120
2	2	317
3	3	222
4	7	100
5	13	312
6	19	50
7	7	120
8	19	400
9	7	230

Output:

name	travelled_distance
Elvis	450
Lee	450
Bob	317
Jonathan	312
Alex	222
Alice	120
Donald	0

Explanation:

Elvis and Lee travelled 450 miles, Elvis is the top traveller as his name is alphabetically smaller than Lee.

Bob, Jonathan, Alex, and Alice have only one ride and we just order them by the total distances of the ride.

Donald did not have any rides, the distance travelled by him is 0.

```
524 •   select user_id, name, sum(distance) as distance_travelled from rides  
525     join users on users.id=rides.user_id  
526     group by 1,2  
527     order by 3 desc,2 asc  
528  
529
```

Result Grid			
	user_id	name	distance_travelled
▶	19	ELVIS	450
	7	LEE	450
	2	BOB	317
	13	JONATHON	312
	3	ALEX	222
	1	ALICE	120

Question 34 Same as Question 26

Question 35

Table: Movies

Column Name	Type
movie_id	int
title	varchar

movie\_id is the primary key for this table.

The title is the name of the movie.

Table: Users

Column Name	Type
user_id	int
name	varchar

user\_id is the primary key for this table.

Table: MovieRating

Column Name	Type
movie_id	int
user_id	int
rating	int
created_at	date

(movie\_id, user\_id) is the primary key for this table.

This table contains the rating of a movie by a user in their review.

created\_at is the user's review date.

Write an SQL query to:

- Find the name of the user who has rated the greatest number of movies. In case of a tie, return the lexicographically smaller user name.
- Find the movie name with the highest average rating in February 2020. In case of a tie, return the lexicographically smaller movie name.

The query result format is in the following example.

-- Find the name of the user who has rated the greatest number of movies

```
584 •   select users.user_id, name, count(*) as rating from movie_rating  
585     join users on users.user_id=movie_rating.user_id  
586     group by 1,2  
587     order by 3 desc ,2 asc  
588     limit 1  
589
```

Result Grid | Filter Rows: \_\_\_\_\_ | Export: | Wrap Cell Content: | Fetch rows:

	user_id	name	rating
▶	1	DANIEL	3

-- Find the movie name with the highest average rating in February 2020.

```
590 •   select title, avg(rating) as avg_rating from movie_rating  
591     join movies on movies.movie_id=movie_rating.movie_id and extract(month from created_at)=2  
592     group by 1  
593     order by 2 desc ,1 asc  
594     limit 1
```

Result Grid | Filter Rows: \_\_\_\_\_ | Export: | Wrap Cell Content: | Fetch rows:

	title	avg_rating
▶	FROZEN 2	3.5000

## Question 36

Table: Users

Column Name	Type
id	int
name	varchar

id is the primary key for this table.  
name is the name of the user.

Table: Rides

Column Name	Type
id	int
user_id	int
distance	int

id is the primary key for this table.  
user\_id is the id of the user who travelled the distance "distance".

Write an SQL query to report the distance travelled by each user.  
Return the result table ordered by travelled\_distance in descending order, if two or more users travelled the same distance, order them by their name in ascending order.  
The query result format is in the following example.

Input:

Users table:

id	name
1	Alice
2	Bob
3	Alex
4	Donald
7	Lee
13	Jonathan
19	Elvis

Rides table:

id	user_id	distance
1	1	120
2	2	317
3	3	222
4	7	100
5	13	312
6	19	50

7	7	120
8	19	400
9	7	230

Output:

name	travelled_distance
Elvis	450
Lee	450
Bob	317
Jonathan	312
Alex	222
Alice	120
Donald	0

Explanation:

Elvis and Lee travelled 450 miles, Elvis is the top traveller as his name is alphabetically smaller than Lee.

Bob, Jonathan, Alex, and Alice have only one ride and we just order them by the total distances of the ride.

Donald did not have any rides, the distance travelled by him is 0.

**Answer:Same as Question 33**

**Q37.**

Table: Employees

Column Name	Type
id	int
name	varchar

id is the primary key for this table.

Each row of this table contains the id and the name of an employee in a company.

Table: EmployeeUNI

Column Name	Type
id	int
unique_id	int

(id, unique\_id) is the primary key for this table.

Each row of this table contains the id and the corresponding unique id of an employee in the company.

Write an SQL query to show the unique ID of each user, If a user does not have a unique ID replace just show null.

Return the result table in any order.

The query result format is in the following example.

Input:

Employees table:

id	name
1	Alice
7	Bob
11	Meir
90	Winston
3	Jonathan

EmployeeUNI table:

id	unique_id
3	1
11	2
90	3

Output:

unique_id	name
null	Alice
null	Bob
2	Meir
3	Winston
1	Jonathan

Explanation:

Alice and Bob do not have a unique ID, We will show null instead.

The unique ID of Meir is 2.

The unique ID of Winston is 3.

The unique ID of Jonathan is 1.

**Answer:Same as Question 32**

## Question 38

Table: Departments

Column Name	Type
id	int
name	varchar

id is the primary key of this table.

The table has information about the id of each department of a university.

Table: Students

Column Name	Type
id	int
name	varchar
department_id	int

id is the primary key of this table.

The table has information about the id of each student at a university and the id of the department he/she studies at.

Write an SQL query to find the id and the name of all students who are enrolled in departments that no longer exist.

Return the result table in any order.

The query result format is in the following example.

Input:

Departments table:

id	name
1	Electrical Engineering
7	Computer Engineering
13	Business Administration

Students table:

id	name	department_id
23	Alice	1
1	Bob	7
5	Jennifer	13
2	John	14
4	Jasmine	77
3	Steve	74
6	Luis	1
8	Jonathan	7
7	Daiana	33
11	Madelynn	1

```
635
636 •   select students.id,students.name from students
637      left join departments on departments.id=students.department_id
638      where departments.id is null
639      order by 2
640
```

	id	name
▶	7	DAIANA
	4	JASMINE
	2	JOHN
	3	STEVE

**Q39.**

Table: Calls

Column Name	Type
from_id	int
to_id	int
duration	int

This table does not have a primary key, it may contain duplicates.

This table contains the duration of a phone call between from\_id and to\_id.

from\_id != to\_id

Write an SQL query to report the number of calls and the total call duration between each pair of distinct persons (person1, person2) where person1 < person2.

Return the result table in any order.

The query result format is in the following example.

Input:

Calls table:

from_id	to_id	duration
1	2	59
2	1	11
1	3	20
3	4	100
3	4	200
3	4	200
4	3	499

Output:

person1	person2	call_count	total_duration
1	2	2	70
1	3	1	20
3	4	4	999

```
661 •  select
662   case
663     when from_id<to_id then from_id else to_id end as from_person,
664   case
665     when to_id>from_id then to_id else from_id end as to_person,sum(duration) as call_duration
666   from calls
667   group by 1,2
```

| Result Grid | Filter Rows:  | Export: | Wrap Cell Content:

	from_person	to_person	call_duration
▶	1	2	70
	1	3	20
	3	4	999

**Q40.**

Table: Prices

Column Name	Type
product_id	int
start_date	date
end_date	date
price	int

(product\_id, start\_date, end\_date) is the primary key for this table.

Each row of this table indicates the price of the product\_id in the period from start\_date to end\_date.  
For each product\_id there will be no two overlapping periods. That means there will be no two intersecting periods for the same product\_id.

Table: UnitsSold

Column Name	Type
product_id	int
purchase_date	date
units	int

There is no primary key for this table, it may contain duplicates.

Each row of this table indicates the date, units, and product\_id of each product sold.

Write an SQL query to find the average selling price for each product. average\_price should be rounded to 2 decimal places.

Return the result table in any order.

The query result format is in the following example.

Input:

Prices table:

product_id	start_date	end_date	price
1	2019-02-17	2019-02-28	5
1	2019-03-01	2019-03-22	20
2	2019-02-01	2019-02-20	15
2	2019-02-21	2019-03-31	30

UnitsSold table:

product_id	purchase_date	units
1	2019-02-25	100
1	2019-03-01	15
2	2019-02-10	200
2	2019-03-22	30

Output:

product_id	average_price
1	6.96
2	16.96

Explanation:

Average selling price = Total Price of Product / Number of products sold.

Average selling price for product 1 =  $((100 * 5) + (15 * 20)) / 115 = 6.96$

Average selling price for product 2 =  $((200 * 15) + (30 * 30)) / 230 = 16.96$

**Answers: SAME AS Question 23**

**Q41.**

Table: Warehouse

Column Name	Type
name	varchar
product_id	int
units	int

(name, product\_id) is the primary key for this table.

Each row of this table contains the information of the products in each warehouse.

Table: Products

Column Name	Type
product_id	int
product_name	varchar
Width	int
Length	int
Height	int

product\_id is the primary key for this table.

Each row of this table contains information about the product dimensions (Width, Length, and Height) in feet of each product.

Write an SQL query to report the number of cubic feet of volume the inventory occupies in each warehouse.

Return the result table in any order.

The query result format is in the following example.

Input:

Warehouse table:

name	product_id	units
LCHouse1	1	1
LCHouse1	2	10
LCHouse1	3	5
LCHouse2	1	2
LCHouse2	2	2
LCHouse3	4	1

Products table:

product_id	product_name	Width	Length	Height
1	LC-TV	5	50	40
2	LC-KeyChain	5	5	5
3	LC-Phone	2	10	10
4	LC-T-Shirt	4	10	20

Output:

warehouse_name	volume
LCHouse1	12250
LCHouse2	20250
LCHouse3	800

```
715 • select name,sum(units*width*length*height) as volume from warehouse  
716     join products on products.product_id=warehouse.product_id  
717     group by 1  
718  
719  
720
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	name	volume		
▶	LCHOUSE1	12250		
	LCHOUSE2	20250		
	LCHOUSE3	800		

## Question 42

Table: Sales

Column Name	Type
sale_date	date
fruit	enum
sold_num	int

(sale\_date, fruit) is the primary key for this table.

This table contains the sales of "apples" and "oranges" sold each day.

Write an SQL query to report the difference between the number of apples and oranges sold each day.  
Return the result table ordered by sale\_date.

The query result format is in the following example.

Input:

Sales table:

sale_date	fruit	sold_num
2020-05-01	apples	10
2020-05-01	oranges	8
2020-05-02	apples	15
2020-05-02	oranges	15
2020-05-03	apples	20
2020-05-03	oranges	0
2020-05-04	apples	15
2020-05-04	oranges	16

Output:

sale_date	diff
2020-05-01	2
2020-05-02	0
2020-05-03	20
2020-05-04	-1

Explanation:

Day 2020-05-01, 10 apples and 8 oranges were sold (Difference  $10 - 8 = 2$ ).

Day 2020-05-02, 15 apples and 15 oranges were sold (Difference  $15 - 15 = 0$ ).

Day 2020-05-03, 20 apples and 0 oranges were sold (Difference  $20 - 0 = 20$ ).

Day 2020-05-04, 15 apples and 16 oranges were sold (Difference  $15 - 16 = -1$ ).

```
742 • Ⓛ select sale_date,diff from (
743   Ⓛ select *,sold-lead(sold) over(partition by sale_date order by fruit asc ) as diff from (
744     select sale_date,fruit,sum(sold_num)  as sold from sales
745     group by 1,2)b
746   )b
747   where diff is not null
748
749
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
sale_date	diff			
2020-05-01	2			
2020-05-02	0			
2020-05-03	20			
2020-05-04	-1			

### Question 43

Table: Activity

Column Name	Type
player_id	int
device_id	int
event_date	date
games_played	int

(player\_id, event\_date) is the primary key of this table.

This table shows the activity of players of some games.

Each row is a record of a player who logged in and played a number of games (possibly 0) before logging out on someday using some device.

Write an SQL query to report the fraction of players that logged in again on the day after the day they first logged in, rounded to 2 decimal places. In other words, you need to count the number of players that logged in for at least two consecutive days starting from their first login date, then divide that number by the total number of players.

The query result format is in the following example.

Input:

Activity table:

player_id	device_id	event_date	games_played
1	2	2016-03-01	5
1	2	2016-03-02	6
2	3	2017-06-25	1
3	1	2016-03-02	0
3	4	2018-07-03	5

Output:

fraction
0.33

Explanation:

Only the player with id 1 logged back in after the first day he had logged in so the answer is  $1/3 = 0.33$

```
775 •   select
776     round(count(case when date_diff = 1 then player_id end)/count(distinct player_id),2) as fraction
777   from (
778     select player_id,event_date,datediff(lead(event_date) over(partition by player_id),event_date)
779       as date_diff
780     from activity
781   )b
782 
```

Result Grid | Filter Rows: \_\_\_\_\_ | Export: | Wrap Cell Content:

	fraction
▶	0.33

#### Question 44

Table: Employee

Column Name	Type
id	int
name	varchar
department	varchar
managerId	int

id is the primary key column for this table.

Each row of this table indicates the name of an employee, their department, and the id of their manager.

If managerId is null, then the employee does not have a manager.

No employee will be the manager of themself.

Write an SQL query to report the managers with at least five direct reports.

Return the result table in any order.

The query result format is in the following example.

Input:

Employee table:

id	name	department	managerId
101	John	A	None
102	Dan	A	101
103	James	A	101
104	Amy	A	101
105	Anne	A	101
106	Ron	B	101

Output:

name
John

```
809  
810 •      select e.name as manager,count(e.name) as direct_reportee  
811          from employee  
812      left join employee as e on e.id=employee.manager_id  
813      group by 1  
814      order by 2 desc  
815      limit 1  
816
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:	Fetch rows
	manager	direct_reportee			
▶	JOHN	5			

#### Question 45

Table: Student

Column Name	Type
student_id	int
student_name	varchar
gender	varchar
dept_id	int

student\_id is the primary key column for this table.

dept\_id is a foreign key to dept\_id in the Department tables.

Each row of this table indicates the name of a student, their gender, and the id of their department.

Table: Department

Column Name	Type
dept_id	int
dept_name	varchar

dept\_id is the primary key column for this table.

Each row of this table contains the id and the name of a department.

Write an SQL query to report the respective department name and number of students majoring in each department for all departments in the Department table (even ones with no current students). Return the result table ordered by student\_number in descending order. In case of a tie, order them by dept\_name alphabetically.

The query result format is in the following example.

Input:

Student table:

student_id	student_name	gender	dept_id
1	Jack	M	1
2	Jane	F	1
3	Mark	M	2

Department table:

dept_id	dept_name
1	Engineering
2	Science
3	Law

```
853 • select department_name,count(case when student.dept_id is not null then 1 end) as students_majoring  
854   from department  
855   left join student on student.dept_id=department.dept_id  
856   group by 1  
857  
858  
859  
860
```

Result Grid		Filter Rows:	Export:	Wrap Cell Content:
	department_name	students_majoring		
▶	ENGINEERING	2		
	SCIENCE	1		
	LAW	0		

**Q46.**

Table: Customer

Column Name	Type
customer_id	int
product_key	int

There is no primary key for this table. It may contain duplicates.  
product\_key is a foreign key to the Product table.

Table: Product

Column Name	Type
product_key	int

product\_key is the primary key column for this table.

Write an SQL query to report the customer ids from the Customer table that bought all the products in the Product table.

Return the result table in any order.

The query result format is in the following example.

Input:

Customer table:

customer_id	product_key
1	5
2	6
3	5
3	6
1	6

Product table:

product_key
5
6

Output:

customer_id
1
3

```
885 • with products as (
886     select *,count(*) over () as total_products from product
887 )
888     select customer_id,total_products,count(distinct case when customer.product_key is not null then customer.product_key end)
889     distinct_products
890     from products
891     left join customer on customer.product_key=products.product_key
892     group by 1,2
893     having distinct_products=total_products
894
```

| Result Grid | Filter Rows: | Export: | Wrap Cell Content:

	customer_id	total_products	distinct_products
▶	1	2	2
	3	2	2

**Q47.**

Table: Project

Column Name	Type
project_id	int
employee_id	int

(project\_id, employee\_id) is the primary key of this table.

employee\_id is a foreign key to the Employee table.

Each row of this table indicates that the employee with employee\_id is working on the project with project\_id.

Table: Employee

Column Name	Type
employee_id	int
name	varchar
experience_years	int

employee\_id is the primary key of this table.

Each row of this table contains information about one employee.

Write an SQL query that reports the most experienced employees in each project. In case of a tie, report all employees with the maximum number of experience years.

Return the result table in any order.

The query result format is in the following example.

Input:

Project table:

project_id	employee_id
1	1
1	2
1	3
2	1
2	4

Employee table:

employee_id	name	experience_years
1	Khaled	3
2	Ali	2
3	John	3
4	Doe	2

Output:

project_id	employee_id
1	1
1	3
2	1

Explanation:

Both employees with id 1 and 3 have the most experience among the employees of the first project.

For the second project, the employee with id 1 has the most experience.

```
934 • ⚏ select * from (
935   select project_id, name, experience_years, dense_rank() over (partition by project.project_id order by experience_years desc) as ranking
936   from project
937   left join employee on employee.employee_id=project.employee_id
938 )b
939 where ranking=1
940
941
942
943
```

Result Grid | Filter Rows: \_\_\_\_\_ | Export: \_\_\_\_\_ | Wrap Cell Content: \_\_\_\_\_

project_id	name	experience_years	ranking
1	KHALED	3	1
1	JOHN	3	1
2	KHALED	3	1

**Q48.**

Table: Books

Column Name	Type
book_id	int
name	varchar
available_from	date

book\_id is the primary key of this table.

Table: Orders

Column Name	Type
order_id	int
book_id	int
quantity	int
dispatch_date	date

order\_id is the primary key of this table.

book\_id is a foreign key to the Books table.

Write an SQL query that reports the books that have sold less than 10 copies in the last year, excluding books that have been available for less than one month from today. Assume today is 2019-06-23.

Return the result table in any order.

The query result format is in the following example.

Input:

Books table:

book_id	name	available_from
1	"Kalila And Demna"	2010-01-01
2	"28 Letters"	2012-05-12
3	"The Hobbit"	2019-06-10
4	"13 Reasons Why"	2019-06-01
5	"The Hunger Games"	2008-09-21

```
987 •   select name,books.book_id,sum(case when extract(year from '2019-06-23')-extract(year from dispatch_date)=1 then quantity else 0 end) as  
988     total_books_sold_last_year  
989   from orders  
990   left join books on books.book_id=orders.book_id  
991   where datediff('2019-06-23',available_from)>30  
991   group by 1,2  
992   having total_books_sold_last_year<=10  
993
```

Result Grid			
	name	book_id	total_books_sold_last_year
▶	Kalila And Demna	1	3
	The Hunger Games	5	0

**Q49.**

Table: Enrollments

Column Name	Type
student_id	int
course_id	int
grade	int

(student\_id, course\_id) is the primary key of this table.

Write a SQL query to find the highest grade with its corresponding course for each student. In case of a tie, you should find the course with the smallest course\_id.

Return the result table ordered by student\_id in ascending order.

The query result format is in the following example.

**Input:**

Enrollments table:

student_id	course_id	grade
2	2	95
2	3	95
1	1	90
1	2	99
3	1	80
3	2	75
3	3	82

**Output:**

student_id	course_id	grade
1	2	99
2	2	95
3	3	82

```
1020
1021 • Ⓜ select * from (
1022   select *,dense_rank() over(partition by student_id order by grade desc,course_id asc) as ranking
1023   from enrollments
1024 )
1025   where ranking=1
1026
1027
```

| Result Grid | Filter Rows:  | Export: Wrap Cell Content:

	student_id	course_id	grade	ranking
▶	1	2	99	1
	2	2	95	1
	3	3	82	1

## Question 50

Table: Teams

Column Name	Type
team_id	int
team_name	varchar

team\_id is the primary key of this table.

Each row of this table represents a single football team.

Table: Matches

Column Name	Type
match_id	int
host_team	int
guest_team	int
host_goals	int
guest_goals	int

match\_id is the primary key of this table.

Each row is a record of a finished match between two different teams.

Teams host\_team and guest\_team are represented by their IDs in the Teams table (team\_id), and they scored host\_goals and guest\_goals goals, respectively.

The winner in each group is the player who scored the maximum total points within the group. In the case of a tie, the lowest player\_id wins.

Write an SQL query to find the winner in each group.

Return the result table in any order.

The query result format is in the following example.

Input:

Players table:

player_id	group_id
15	1
25	1
30	1
45	1
10	2
35	2
50	2
20	3
40	3

Matches table:

match_id	first_player	second_player	first_score	second_score

1	15	45	3	0
2	30	25	1	2
3	30	15	2	0
4	40	20	5	2
5	35	50	1	1

Output:

group_id	player_id
1	15
2	35
3	40

### Method 1

```

with base as (
select matches.*,players.group_id as first_group,p.group_id as second_group
from matches
left join players on players.player_id=matches.first_player
left join players as p on p.player_id=matches.second_player
)

select * from (
select *,row_number() over (partition by player_group order by total_goals
desc) as ranking
from (
select player_group,player,sum(goals) as total_goals from (
(select first_group as player_group,first_player as
player,sum(first_player_goals) as goals from base
group by 1,2)
union all
(select second_group as player_group,second_player as
player,sum(second_player_goals) as goals from base
group by 1,2)
)b
group by 1,2
)b

```

```
)b  
where ranking=1
```

## Method 2

```
1093 • Ⓜ select * from (
1094     select player_id,group_id,sum(if(first_player=player_id,first_player_goals,second_player_goals)) as scores,row_number() over (partition by
1095         group_id order by sum(if(first_player=player_id,first_player_goals,second_player_goals)) desc) as ranking
1096     from matches
1097     join players on players.player_id=matches.first_player or players.player_id=matches.second_player
1098     group by 1,2
1099 )b
1100 where ranking=1
1100
```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

player_id	group_id	scores	ranking
15	1	3	1
35	2	1	1
40	3	5	1

Result Grid | Form Editor

